

Phase 3 Project

1. Business understanding

Project Goal

To predict whether there is a pattern of customers who will ("soon") stop doing business with SyriaTel, a telecommunications company.

The audience are the telecom business staff whose interest is reducing how much money is lost because of customers who don't want to stay with the company very long.

Objectives

To determine if there is a predictive pattern of customers who will ("soon") stop doing business with SyriaTel.

2. Data understanding

2.1 Loading the data

```
In [1]: # importing the necessary Libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from scipy import stats  
import warnings  
warnings.filterwarnings("ignore")  
import statsmodels.api as sm  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import mean_squared_error, r2_score  
from sklearn.metrics import confusion_matrix  
from sklearn.preprocessing import LabelEncoder
```

```
In [2]: # Importing the data
df = pd.read_csv('bigml_csv.csv', index_col = 0)
df.head() # returns 3333 entries and 20 columns
```

Out[2]:

	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	t
state										
KS	128	415	382-4657	no	yes	25	265.1	110	45.07	1
OH	107	415	371-7191	no	yes	26	161.6	123	27.47	1
NJ	137	415	358-1921	no	no	0	243.4	114	41.38	1
OH	84	408	375-9999	yes	no	0	299.4	71	50.90	1
OK	75	415	330-6626	yes	no	0	166.7	113	28.34	1

```
In [3]: # more information about the data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3333 entries, KS to TN
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   account length                        3333 non-null   int64
1   area code                            3333 non-null   int64
2   phone number                         3333 non-null   object
3   international plan                   3333 non-null   object
4   voice mail plan                      3333 non-null   object
5   number vmail messages                3333 non-null   int64
6   total day minutes                    3333 non-null   float64
7   total day calls                      3333 non-null   int64
8   total day charge                     3333 non-null   float64
9   total eve minutes                    3333 non-null   float64
10  total eve calls                      3333 non-null   int64
11  total eve charge                     3333 non-null   float64
12  total night minutes                  3333 non-null   float64
13  total night calls                    3333 non-null   int64
14  total night charge                   3333 non-null   float64
15  total intl minutes                   3333 non-null   float64
16  total intl calls                     3333 non-null   int64
17  total intl charge                    3333 non-null   float64
18  customer service calls               3333 non-null   int64
19  churn                               3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(3)
memory usage: 524.0+ KB
```

3. Data Preparation

```
In [4]: ▶ # format the naming of the columns # Replace the spaces with '_'
df.rename(columns=lambda x: x.replace(' ', '_'), inplace=True)
```

```
In [5]: ▶ # check the formatted column names
print(df.columns)
```

```
Index(['account_length', 'area_code', 'phone_number', 'international_pla
n',
      'voice_mail_plan', 'number_vmail_messages', 'total_day_minutes',
      'total_day_calls', 'total_day_charge', 'total_eve_minutes',
      'total_eve_calls', 'total_eve_charge', 'total_night_minutes',
      'total_night_calls', 'total_night_charge', 'total_intl_minutes',
      'total_intl_calls', 'total_intl_charge', 'customer_service_calls',
      'churn'],
      dtype='object')
```

3.1. Check for missing values

```
In [6]: ▶ # checking for missing values
df.isnull().sum()
```

```
Out[6]: account_length      0
area_code                  0
phone_number               0
international_plan         0
voice_mail_plan            0
number_vmail_messages      0
total_day_minutes          0
total_day_calls            0
total_day_charge           0
total_eve_minutes          0
total_eve_calls            0
total_eve_charge           0
total_night_minutes        0
total_night_calls          0
total_night_charge         0
total_intl_minutes         0
total_intl_calls           0
total_intl_charge          0
customer_service_calls     0
churn                     0
dtype: int64
```

There are no missing values

```
In [7]: ▶ # Format the columns to remove any spaces from the column lables
df.columns = df.columns.str.strip()
```

```
In [8]: ▶ # print the column names
print(df.columns)
```

```
Index(['account_length', 'area_code', 'phone_number', 'international_pla
n',
      'voice_mail_plan', 'number_vmail_messages', 'total_day_minutes',
      'total_day_calls', 'total_day_charge', 'total_eve_minutes',
      'total_eve_calls', 'total_eve_charge', 'total_night_minutes',
      'total_night_calls', 'total_night_charge', 'total_intl_minutes',
      'total_intl_calls', 'total_intl_charge', 'customer_service_calls',
      'churn'],
      dtype='object')
```

```
In [9]: ▶ # print the unique values of area code
df['area_code'].unique()
```

```
Out[9]: array([415, 408, 510], dtype=int64)
```

3.2. Defining the features (X) and target (y)

```
In [10]: ▶ # Defining the features (X) and target (y)
X_features = df.drop('churn', axis=1)
y_target = df['churn']
```

```
In [11]: ▶ #convert categorical variables to numerical
# drop-first = True
X_numerical = pd.get_dummies(X_features, drop_first = True)
X_numerical.head(5)
```

```
Out[11]:
```

	account_length	area_code	number_vmail_messages	total_day_minutes	total_day_cal
state					
KS	128	415	25	265.1	1
OH	107	415	26	161.6	1
NJ	137	415	0	243.4	1
OH	84	408	0	299.4	1
OK	75	415	0	166.7	1

5 rows × 3350 columns



```
In [12]: ▶ # Encode 'churn'
# print y_encode
y_encode = y_target.astype(int)
y_encode
```

```
Out[12]: state
KS      0
OH      0
NJ      0
OH      0
OK      0
..
AZ      0
WV      0
RI      0
CT      0
TN      0
Name: churn, Length: 3333, dtype: int32
```

```
In [13]: ▶ #encode area code
df['area_code'] = LabelEncoder().fit_transform(df['area_code'])
df['area_code']
```

```
Out[13]: state
KS      1
OH      1
NJ      1
OH      0
OK      1
..
AZ      1
WV      1
RI      2
CT      2
TN      1
Name: area_code, Length: 3333, dtype: int64
```

3.3 Selecting the features for modelling

```
In [14]: ▶ # Selecting the features for medelling
X = X_numerical[['account_length', 'area_code', 'total_day_charge', 'total
y = y_encode
```

In [15]: `# print X`
`print(X)`

state	account_length	area_code	total_day_charge	total_eve_charge	\
KS	128	415	45.07	16.78	
OH	107	415	27.47	16.62	
NJ	137	415	41.38	10.30	
OH	84	408	50.90	5.26	
OK	75	415	28.34	12.61	
...	
AZ	192	415	26.55	18.32	
WV	68	415	39.29	13.04	
RI	28	510	30.74	24.55	
CT	184	510	36.35	13.57	
TN	74	415	39.85	22.60	

state	total_night_charge	total_intl_charge
KS	11.01	2.70
OH	11.45	3.70
NJ	7.32	3.29
OH	8.86	1.78
OK	8.41	2.73
...
AZ	12.56	2.67
WV	8.61	2.59
RI	8.64	3.81
CT	6.26	1.35
TN	10.86	3.70

[3333 rows x 6 columns]

In [16]: `#print y`
`y`

Out[16]:

state	
KS	0
OH	0
NJ	0
OH	0
OK	0
..	
AZ	0
WV	0
RI	0
CT	0
TN	0

Name: churn, Length: 3333, dtype: int32

4. Modelling

4.1 Modelling Using statsmodels

```
In [17]: ▶ # fit the baseline Logistic Regression model using statsmodels  
# Add a constant (intercept) term to the features  
X_constant = sm.add_constant(X)  
  
# Create and fit the Logistic Regression model  
logistic_model = sm.Logit(y, X_constant)  
results = logistic_model.fit()
```

```
Optimization terminated successfully.  
    Current function value: 0.384108  
    Iterations 7
```

```
In [18]: # Print summary
print(results.summary())
```

```

                                Logit Regression Results
=====
Dep. Variable:                  churn    No. Observations:
3333                                     3333
Model:                          Logit    Df Residuals:
3326                                     3326
Method:                          MLE    Df Model:
6                                     6
Date:                            Sat, 10 May 2025    Pseudo R-squ.:          0.
07172                                     0.07172
Time:                            17:59:59    Log-Likelihood:          -1
280.2                                     -1280.2
converged:                        True    LL-Null:          -1
379.1                                     -1379.1
Covariance Type:                nonrobust    LLR p-value:          5.50
6e-40
=====
=====
                                coef    std err          z      P>|z|      [0.02
5      0.975]
-----
const                -6.8268      0.696     -9.806      0.000     -8.19
1      -5.462
account_length        0.0012      0.001      0.954      0.340     -0.00
1      0.004
area_code             0.0007      0.001      0.556      0.578     -0.00
2      0.003
total_day_charge      0.0675      0.006     11.632      0.000      0.05
6      0.079
total_eve_charge      0.0662      0.012      5.483      0.000      0.04
3      0.090
total_night_charge    0.0505      0.022      2.265      0.023      0.00
7      0.094
total_intl_charge     0.2935      0.069      4.262      0.000      0.15
9      0.428
=====
=====

```

Pseudo R-Squared of 0.07172 indicated that the model explains only 7.17% of the variability

```
In [19]: # Printing the coefficients
print(results.params)
```

```

const                -6.826829
account_length        0.001220
area_code             0.000664
total_day_charge      0.067466
total_eve_charge      0.066206
total_night_charge    0.050466
total_intl_charge     0.293521
dtype: float64

```



The coefficients indicate that `total_day_charge`, `total_eve_charge`, `total_night_charge`, `total_intl_charge` are significant and influence whether one churns or not.

A unit increase in international charge increases the chance of churn by 29%.

4.2. Modelling using scikit-learn

```
In [20]: ▶ # import the necessary library
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score
```

```
In [21]: ▶ # Select the features (X) and target (y)
X = X_numerical[['account_length', 'area_code', 'total_day_charge', 'total
y = y_encode
```



```
In [22]: ▶ # train and split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

```
In [23]: ▶ #view the lengths of the results
print(len(X_train), len(X_test), len(y_train), len(y_test))
```

2666 667 2666 667

```
In [24]: ▶ # Train and fit model
model2 = LogisticRegression()
results2 = model2.fit(X_train, y_train)
```

In [25]: `print(results.summary())`

```

                                Logit Regression Results
=====
=====
Dep. Variable:                  churn    No. Observations:
3333
Model:                          Logit    Df Residuals:
3326
Method:                         MLE     Df Model:
6
Date:                          Sat, 10 May 2025    Pseudo R-squ.:          0.
07172
Time:                          18:00:18    Log-Likelihood:          -1
280.2
converged:                      True     LL-Null:          -1
379.1
Covariance Type:                nonrobust    LLR p-value:          5.50
6e-40
=====
=====
                                coef    std err          z      P>|z|      [0.02
5      0.975]
-----
-----
const                -6.8268      0.696     -9.806      0.000     -8.19
1      -5.462
account_length      0.0012      0.001      0.954      0.340     -0.00
1      0.004
area_code           0.0007      0.001      0.556      0.578     -0.00
2      0.003
total_day_charge    0.0675      0.006    11.632      0.000      0.05
6      0.079
total_eve_charge    0.0662      0.012      5.483      0.000      0.04
3      0.090
total_night_charge  0.0505      0.022      2.265      0.023      0.00
7      0.094
total_intl_charge   0.2935      0.069      4.262      0.000      0.15
9      0.428
=====
=====


```

In [26]: `# Make predictions on the test data`
`y_pred = model2.predict(X_test)`

In [27]: `#plotting coefficients vs features`

```
coefs = pd.Series(model2.coef_[0], index=X.columns).sort_values()

plt.figure(figsize=(8, 5))
coefs.plot(kind='barh')
plt.title('Logistic Regression Feature Importance (Coefficients)')
plt.xlabel('Coefficient Value')
plt.grid(True)
plt.show()
```



In [28]: `coefs`

```
Out[28]: account_length      0.001590
         area_code          0.002292
         total_night_charge  0.034522
         total_eve_charge    0.055797
         total_day_charge    0.062950
         total_intl_charge   0.316269
         dtype: float64
```

The most important feature is the total_intl_charge

A unit increase in international charge increases the chance of churn by 31%.

In [29]: `#Evaluation using mse and r squared`

```
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred) # 0-1: the higher (closer to 1) the better th
print("mse", mse)
print("r2:", r2)
```

```
mse 0.14992503748125938
r2: -0.16677745513067221
```

A negative r2 indicated that the model is performing poorly

```
In [30]: ▶ # Evaluate the model on training data
train_pred = model2.predict(X_train)
print("Training Accuracy:", accuracy_score(y_train, train_pred)) # checking
print("Training report")
print(classification_report(y_train, train_pred))

#f1 is a combination of precision and recall
```

Training Accuracy: 0.858589647411853

Training report

	precision	recall	f1-score	support
0	0.86	1.00	0.92	2284
1	1.00	0.01	0.03	382
accuracy			0.86	2666
macro avg	0.93	0.51	0.47	2666
weighted avg	0.88	0.86	0.80	2666

The results indicate that the model accurately predicts whether one will churn or not churn SyriaTel 86% of the time when using the training data

The f1 score for the class 0 is 0.92. However, the f1 score for class 1 is very low 0.02. This is due to the data imbalances, there are more instances for class 0 (2284) compared to class 1(382)

4.3 Confusion Matrix

```
In [32]: ▶ # import library
from sklearn.metrics import confusion_matrix
```

```
In [33]: ▶ # Evaluate on testing data
test_pred = model2.predict(X_test)
print("Testing Accuracy:", accuracy_score(y_test, test_pred)) # checking t
print("Testing report")
print(classification_report(y_test, test_pred))
```

Testing Accuracy: 0.8500749625187406

Testing report

	precision	recall	f1-score	support
0	0.85	1.00	0.92	566
1	1.00	0.01	0.02	101
accuracy			0.85	667
macro avg	0.92	0.50	0.47	667
weighted avg	0.87	0.85	0.78	667

5. Evaluation using ROC and AUC

```
In [34]: ▶ # addressing the data imbalances

from sklearn.metrics import roc_curve, auc

# Scikit-Learn's built in roc_curve method returns the fpr, tpr, and thres

# Calculate the probability scores of each of the datapoints:
y_score = model2.fit(X_train, y_train).decision_function(X_test)

fpr, tpr, thresholds = roc_curve(y_test, y_score)
```

```
In [35]: ▶ roc_auc = auc(fpr, tpr)
roc_auc
```

Out[35]: 0.7116467830528636

```
In [36]: ▶ # Print the AUC
print('AUC: {}'.format(auc(fpr, tpr)))
```

AUC: 0.7116467830528636

The model has a moderate ability of prediction. The model has a 71.16% chance of predictily correctly the negative and positive classes.

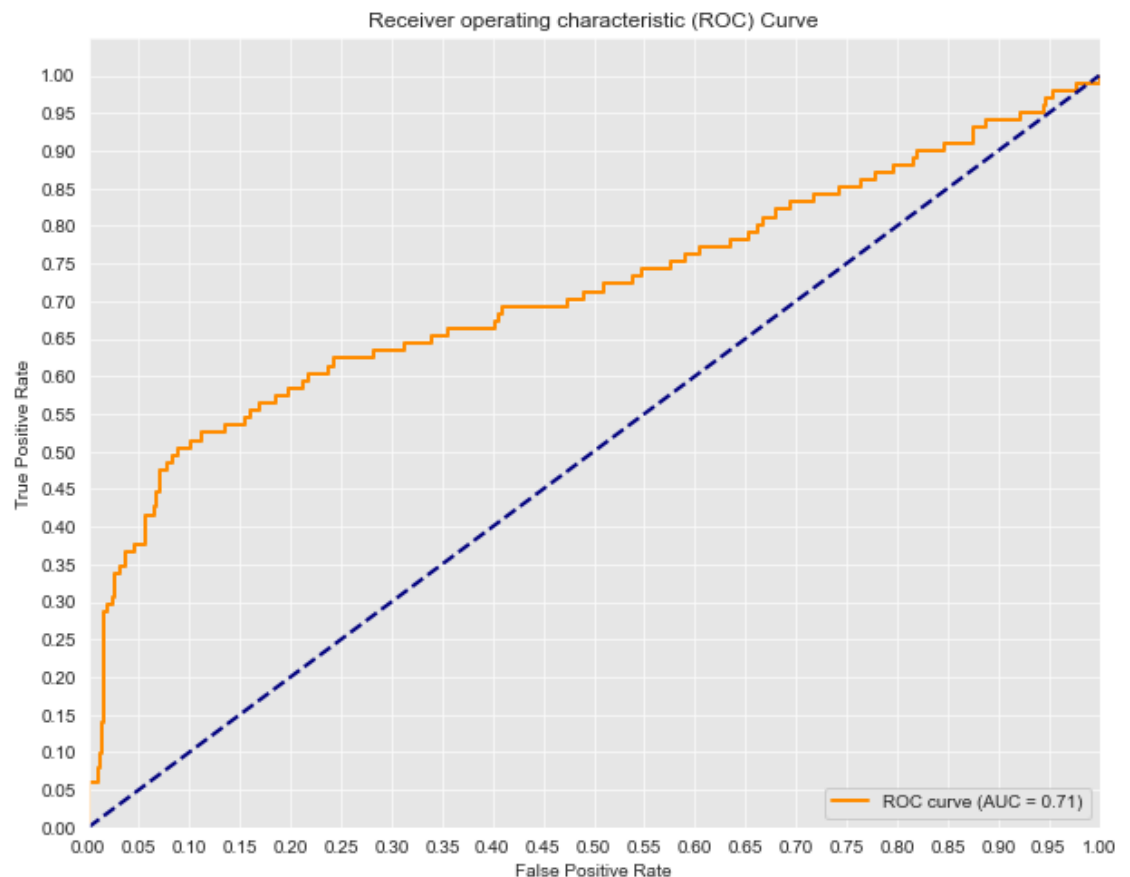
```
In [37]: # Visualizing the ROC and AUC

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# Seaborn's beautiful styling
sns.set_style('darkgrid', {'axes.facecolor': '0.9'})

print('AUC: {}'.format(auc(fpr, tpr)))
plt.figure(figsize=(10, 8))
lw = 2
plt.plot(fpr, tpr, color='darkorange',
         lw=lw, label='ROC curve (AUC = {:.2f})'.format(roc_auc))
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.yticks([i/20.0 for i in range(21)])
plt.xticks([i/20.0 for i in range(21)])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```

AUC: 0.7116467830528636



Scaling the data

```
In [38]: # import the libraries
from sklearn.preprocessing import MinMaxScaler, PolynomialFeatures, StandardScaler

X = X_numerical[['account_length', 'area_code', 'total_day_charge', 'total_eve_charge', 'total_intl_charge']]
y = y_encode

# train and split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

# Standardize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# train and fit the model
model3 = LogisticRegression()
results3 = model3.fit(X_train_scaled, y_train)
results3
```

Out[38]:

▼ LogisticRegression

LogisticRegression()

Hyperparameter Tuning in Decision Trees

```
In [52]: #import the Libraries
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_squared_error
```

```
In [53]: # define the data
X = X_numerical[['account_length', 'total_day_charge', 'total_eve_charge', 'total_intl_charge']]
y = df['churn']
```

```
In [54]: # Train and split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
```

```
In [55]: ▶ # Define the parameter grid to tune the hyperparameters
param_grid = {
    'max_depth': [10, 20, 30, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

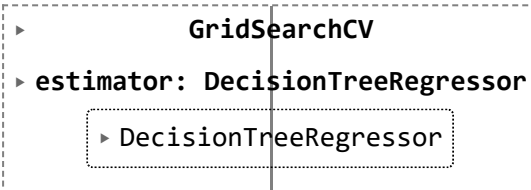
# Instantiating the model
dtree_reg = DecisionTreeRegressor(random_state=1000)

# Perform the grid search
grid_search = GridSearchCV(estimator=dtree_reg, param_grid=param_grid,
                           cv=5, n_jobs=-1, verbose=2, scoring='neg_mean_s
```

```
In [56]: ▶ # Fit the model
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 36 candidates, totalling 180 fits

```
Out[56]:
```



```

    └─ GridSearchCV
      └─ estimator: DecisionTreeRegressor
         └─ DecisionTreeRegressor

```

```
In [57]: ▶ # Get the best estimator from the grid search
best_dtree_reg = grid_search.best_estimator_

# Predict on the test set
y_pred = best_dtree_reg.predict(X_test)

# Calculate Mean Squared Error and RMSE
mse = mean_squared_error(y_test, y_pred)
rmse = mse ** 0.5
```

```
In [58]: ▶ # Get the best hyperparameters
best_params = grid_search.best_params_
print(f"Best parameters: {best_params}")
print(f"Test RMSE: {rmse}")
```

```
Best parameters: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_sp
lit': 10}
Test RMSE: 0.33638802926946276
```

RMSE is 0.336 which is better than the baseline model where the RMSE was negative

Type Markdown and LaTeX: α^2

GitHub repository link: <https://github.com/PamGodia/dsc-phase3-project-pam> (<https://github.com/PamGodia/dsc-phase3-project-pam>)

